



UNIVERSITY OF COPENHAGEN

Recovery: Basic Concepts

Recovery: ARIES, normal operation

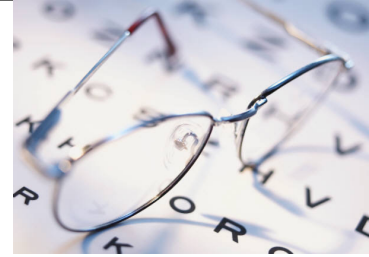
ACS, Dmitriy Traytel

Do-it-yourself-recap: The many faces of atomicity

- **Atomicity** is strong modularity mechanism!
 - Hides that one high-level action is actually made of many sub-actions
- **Before-or-after** atomicity
 - == Isolation
 - Cannot have effects that would only arise by interleaving of parts of transactions
- **What was the meaning of all-or-nothing atomicity?**



What should we learn today?

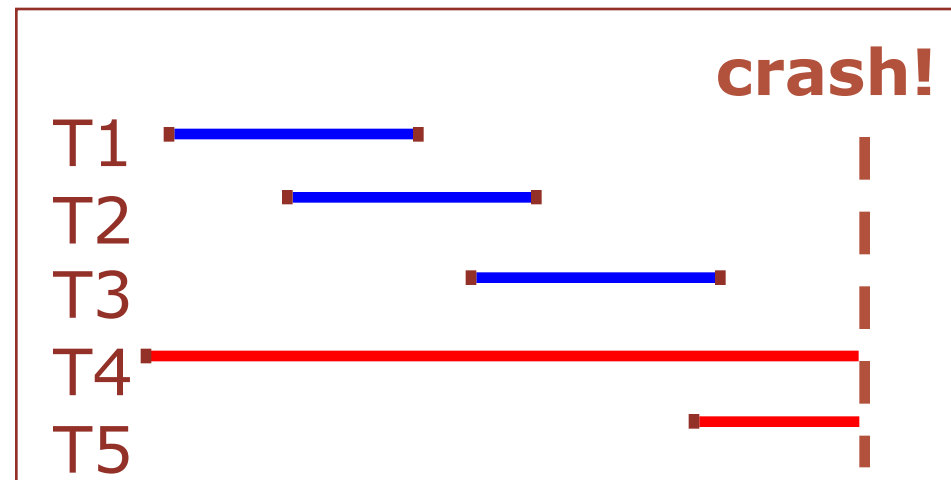


- Explain the concepts of volatile, nonvolatile, and stable storage as well as the main assumptions underlying database recovery
- Predict how force/no-force and steal/no-steal strategies for writes and buffer management influence the need for redo and undo
- Explain the notion of logging and the concept of write-ahead logging
- Predict what portions of the log and database are necessary for recovery based on the recovery equations
- Explain how write-ahead logging is achieved in the ARIES protocol
- Explain the functions of recovery metadata such as the transaction table and the dirty page table
- Interpret the contents of the log resulting from ARIES normal operation

Implementing All-or-Nothing Atomicity

- Atomicity
 - Transactions may abort (“Rollback”).
- Durability
 - What if system stops running? (Causes?)

- ❖ Desired Behavior after system restarts:
- T1, T2 & T3 should be durable.
 - T4 & T5 should be aborted (effects not seen).



Assumptions

- Concurrency control is in effect
 - **Strict 2PL**, in particular
- Updates are happening “in place”
 - i.e. data is overwritten on (deleted from) memory using `READ / WRITE` interface.
 - We will use a two-level memory with buffer and disk
- Types of failures
 - Crash
 - Media failure
- Always fail-stop!



Volatile vs. Nonvolatile vs. Stable Storage

- **Volatile Storage**
 - Lost in the event of a crash
 - Example: main memory
- **Nonvolatile Storage**
 - Not lost on crash, but lost on media failure
 - Example: disk
- **Stable Storage**
 - Never lost (otherwise, that's it 😊)
 - How do you implement this one?



Surviving Crashes: How to handle the Buffer Pool?

- **Force** every write to disk?
 - Poor response time.
 - But provides durability.
- **Steal** buffer-pool frames from uncommitted Xacts?
 - If not, poor throughput.
 - If so, how can we ensure atomicity?

	No Steal	Steal
Force	Trivial (?)	
No Force		Desired



Surviving Crashes: How to handle the Buffer Pool?

Fill in the matrix:
When do you need to UNDO changes?
When do you need to REDO changes?

- **Force** every write to disk?
 - Poor response time.
 - But provides durability.
- **Steal** buffer-pool frames from uncommitted Xacts?
 - If not, poor throughput.
 - If so, how can we ensure atomicity?

	No Steal	Steal
Force	Trivial (?)	
No Force		Desired



More on Steal and Force



More on Steal and Force

- **STEAL** (why enforcing Atomicity is hard)
 - *To steal frame F:* Current page in F (say P) is written to disk; some Xact holds lock on P.
 - What if the Xact with the lock on P aborts?
 - Must remember the old value of P at steal time (to support **UNDO**ing the write to page P).



More on Steal and Force

- **STEAL** (why enforcing Atomicity is hard)
 - *To steal frame F*: Current page in F (say P) is written to disk; some Xact holds lock on P.
 - What if the Xact with the lock on P aborts?
 - Must remember the old value of P at steal time (to support **UNDO**ing the write to page P).
- **NO FORCE** (why enforcing Durability is hard)
 - What if system crashes before a modified page is written to disk?
 - Write as little as possible, in a convenient place, at commit time, to support **REDO**ing modifications.



Undo/Redo vs. Force/Steal

		No Steal	Steal
Force		No Redo No Undo	No Redo Undo
No Force		Redo No Undo	Redo Undo



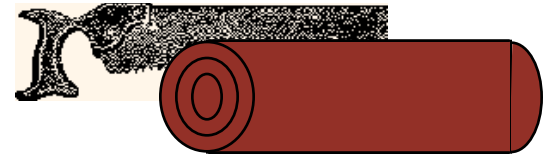
Undo/Redo vs. Force/Steal

	No Steal	Steal
Force	No Redo No Undo	No Redo Undo
No Force	Redo No Undo	Redo Undo

How do we support this option?



Basic Idea: Logging



- Record REDO and UNDO information, for every update, in a *log*.
 - Sequential writes to log (put it on a separate disk).
 - Minimal info (diff) written to log, so multiple updates fit in a single log page.
- Log: An ordered list of REDO/UNDO actions
 - Logical vs. Physical Logging
 - Example physical log record contains:

<XID, pageID, offset, length, old data, new data>

- Good compromise is physiological logging.

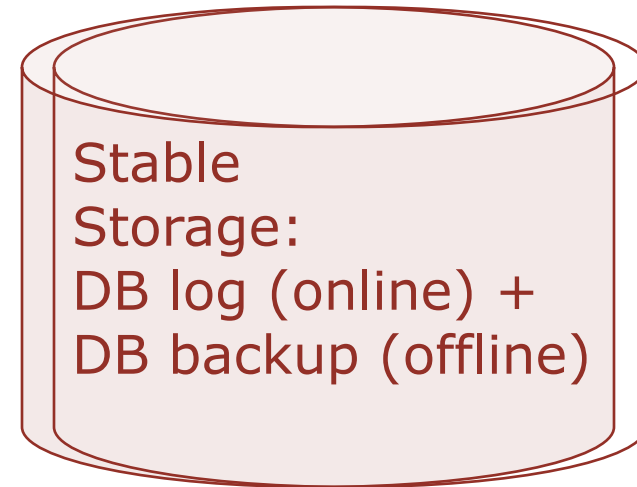



Write-Ahead Logging (WAL)

- Golden Rule: Never modify the only copy!
- The **Write-Ahead Logging** Protocol:
 - 1) Must **force** the **log record** for an update before the corresponding **data page** gets to disk.
 - 2) Must **write all log records** for a Xact before commit.
- #1 guarantees Atomicity.
- #2 guarantees Durability.
- Exactly how is logging (and recovery!) done?
 - We will study the ARIES algorithms.



Recovery Equations



- Crash Recovery: volatile memory lost **Discussion:**
 - Current DB = DB files + DB log \longrightarrow since when?
- Media Recovery: nonvolatile storage lost
 - Current DB = DB backup + DB log \longrightarrow since when?
- We will focus on crash recovery next

Undo/Redo vs. Force/Steal

		No Steal	Steal
Force		No Redo No Undo	No Redo Undo
No Force		Redo No Undo	Redo Undo

Undo/Redo vs. Force/Steal

	No Steal	Steal
Force	No Redo No Undo	No Redo Undo
No Force	Redo No Undo	Redo Undo

How do we support this option?

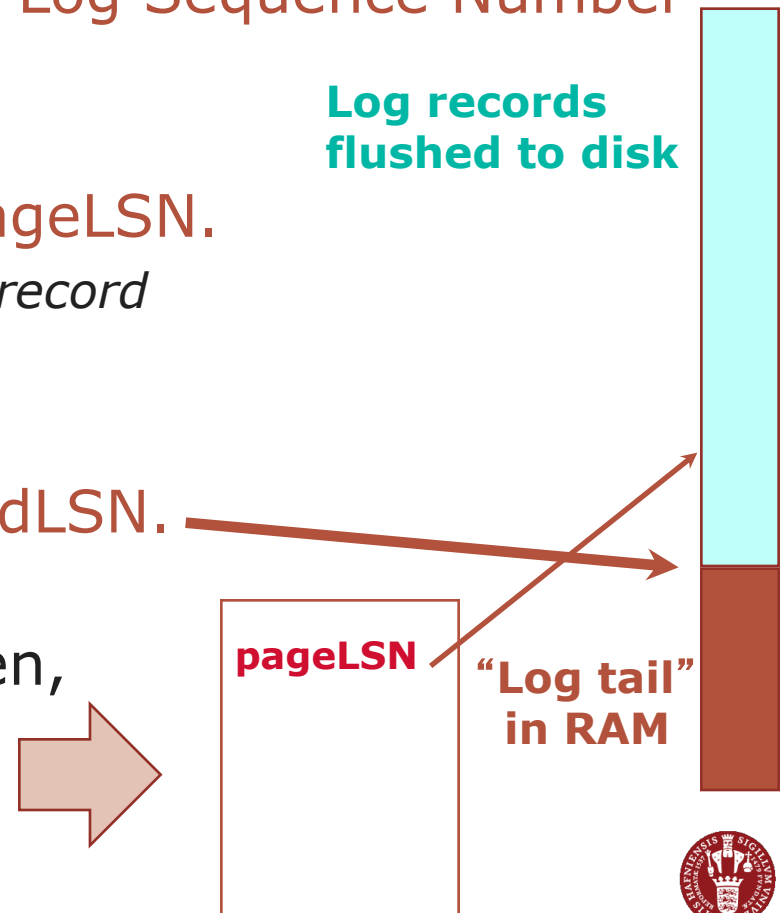


WAL & the Log



- Each log record has a unique **Log Sequence Number (LSN)**.
 - LSNs always increasing.
- Each data page contains a **pageLSN**.
 - The LSN of the most recent *log record* for an update to that page.
- System keeps track of **flushedLSN**.
 - The max LSN flushed so far.
- WAL: *Before* a page is written,
 - $\text{pageLSN} \leq \text{flushedLSN}$

Log records
flushed to disk



Log Records

LogRecord fields:

prevLSN

XID

type

pageID

length

offset

before-image

after-image

**update
records
only**

Possible log record types:

- **Update**
- **Commit**
- **Abort**
- **End** (signifies end of commit or abort)
- **Compensation Log Records (CLRs)**
 - for UNDO actions

Note: Format above simplified; in reality, ARIES uses physiological variant

Source: Ramakrishnan & Gehrke (partial)



Other Log-Related State

- **Transaction Table:**
 - One entry per active Xact.
 - Contains **XID**, **status** (running/committed/aborted), and **lastLSN**.
- **Dirty Page Table:**
 - One entry per dirty page in buffer pool.
 - Contains **recLSN** -- the LSN of the log record which ***first*** caused the page to be dirty.



Normal Execution of an Xact

Keep in Mind:
It must be OK to
crash at **any time**
→ **repeat history!**

- Series of **reads** & **writes**, followed by **commit** or **abort**.
 - We will assume that write is atomic on disk.
 - In practice, additional details to deal with non-atomic writes.
- **Strict 2PL** → concurrency is correctly handled
- **STEAL, NO-FORCE** buffer management,
with **Write-Ahead Logging**.



The Big Picture: What's Stored Where



LogRecords

prevLSN
XID
type
pageID
length
offset
before-image
after-image



Data pages

each
with a
pageLSN

master record



Xact Table

lastLSN
status

Dirty Page Table

recLSN

flushedLSN



Checkpointing

- Periodically, the DBMS creates a checkpoint, to minimize the time taken to recover in the event of a system crash. Write to log:
 - begin_checkpoint** record: Indicates when checkpoint began.
 - end_checkpoint** record: Contains current *Xact table* and *dirty page table*. This is a **'fuzzy checkpoint'**:
 - Other Xacts continue to run; so these tables accurate only as of the time of the **begin_checkpoint** record.
 - No attempt to force dirty pages to disk**; effectiveness of checkpoint limited by oldest unwritten change to a dirty page, **minDirtyPagesLSN**.
 - Use **background process** to flush dirty pages to disk!
 - Store LSN of checkpoint record in a safe place (**master** record).



Transaction Commit

- Write **commit** record to log.
- All log records up to Xact's **lastLSN** are flushed.
 - Guarantees that **flushedLSN** \geq **lastLSN**.
 - Note that log flushes are sequential, synchronous writes to disk.
 - Many log records per log page.
- Commit() returns.
- Write **end** record to log.




Simple Transaction Abort

- For now, consider an explicit abort of a Xact.
 - No crash involved.
- We want to “play back” the log in reverse order, UNDOing updates.
 - Get **lastLSN** of Xact from Xact table.
 - Can follow chain of log records backward via the **prevLSN** field.
 - Before starting UNDO, write an **Abort log record**.
 - For recovering from crash during UNDO!



Abort, cont.



- To perform UNDO, must have a lock on data!
 - Strict 2PL enforces this
- Before restoring old value of a page, write a CLR:
 - You continue logging while you UNDO!!
 - CLR has one extra field: **undonextLSN**
 - Points to the next LSN to undo (i.e. the prevLSN of the record we're currently undoing).
 - CLRs *never* Undone (but they might be Redone when repeating history: guarantees Atomicity!)
- At end of UNDO, write an **end** log record.



Example

- 10 T1 writes P5
- 20 T2 writes P17
- 30 T1 writes P3

P3 written to disk

(pageLSN for page 3 at this time is 30)

- 40 T1 aborts
- 50 CLR T1 P3 (undonextLSN: 10)
- 60 CLR T1 P5 (undonextLSN: NULL)
- 70 End T1



A Longer Example

- 10 T1 writes P3 (prevLSN: NULL)
- 20 T2 writes P4 (prevLSN: NULL)
- 30 T2 writes P5 (prevLSN: 20)
- flushedLSN = 20

P4 gets written to disk (pageLSN for page 4 = 20)

T2 aborts

- 50 Abort T2
- 60 CLR T2 P5 (undoNextLSN = 20), pageLSN(P5)=60

Update P5 in the buffer manager

Flush log up to log record 60

Buffer manager writes P5 to disk.

- 70 CLR T2 P4 (undoNextLSN = NULL), pageLSN(P4)=70

Update P4 in the buffer manager

- 80 End T2
- 90 T1 commits

Flush log up to log record 90, then the commit(T1) returns

Discussion: Does
this example
make sense?
Can you explain
to your
colleague what
happened?

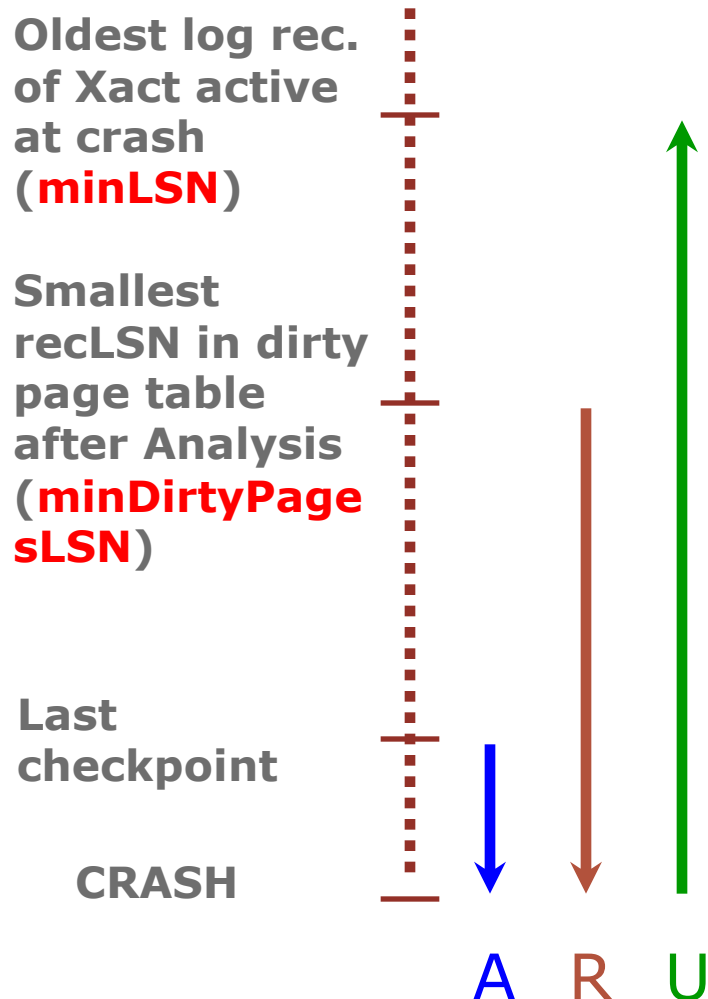


Questions so far?



Crash Recovery: Big Picture

Keep in Mind:
It must be OK to
crash at **any time**
(including during
recovery)



- ❖ Start from a **checkpoint** (found via **master** record).
- ❖ Three phases. Need to:
 - Figure out which Xacts committed since checkpoint, which failed (**Analysis**).
 - **REDO** *all* actions.
 - ◆ **Repeat History**
 - **UNDO** effects of failed Xacts.

Recovery: The Analysis Phase

- Reconstruct state at checkpoint.
 - via **end_checkpoint** record.
- Scan log forward from checkpoint.
 - **End** record: Remove Xact from Xact table.
 - **Other records**: Add Xact to Xact table, set **lastLSN=LSN**, change Xact status on **commit**.
 - **Update** record: If P not in Dirty Page Table,
 - Add P to Dirty Page Table, set its **recLSN=LSN**.



Recovery: The REDO Phase

- We *repeat History* to reconstruct state at crash:
 - Reapply *all* updates (even of aborted Xacts!), redo CLR.
- Scan forward from log record containing the smallest **recLSN** in Dirty Page Table. For each CLR or update log record with **LSN**, REDO the action unless:
 - Affected page is not in the Dirty Page Table, or
 - Affected page is in Dirty Page Table, but has **recLSN** > **LSN**, or
 - **pageLSN** (in DB) >= **LSN**.
- To **REDO** an action:
 - Reapply logged action.
 - Set **pageLSN** to **LSN**. No additional logging! (**Why?**)



Recovery: The UNDO Phase

$ToUndo = \{ lsn \mid lsn \text{ a lastLSN of a "loser" Xact} \}$

Repeat:

- Choose largest LSN among ToUndo.
- If this LSN is a CLR and $undonextLSN == NULL$
 - Write an End record for this Xact.
- If this LSN is a CLR, and $undonextLSN \neq NULL$
 - Add $undonextLSN$ to ToUndo
- Else this LSN is an update. Undo the update, write a CLR, add $prevLSN$ to ToUndo.

Until ToUndo is empty.



Example of Recovery



Xact Table

lastLSN

status

Dirty Page Table

recLSN

flushedLSN

ToUndo

LSN	LOG
00	begin_checkpoint
05	end_checkpoint
10	update: T1 writes P5
20	update T2 writes P3
30	T1 abort
40	CLR: Undo T1 LSN 10
45	T1 End
50	update: T3 writes P1
60	update: T2 writes P5
X CRASH, RESTART	

prevLSNs

Example: Crash During Restart!

Discussion:
Explain to your
colleague!



Xact Table

lastLSN
status

Dirty Page Table

recLSN

flushedLSN

ToUndo

LSN	LOG
00,05	begin_checkpoint, end_checkpoint
10	update: T1 writes P5
20	update T2 writes P3
30	T1 abort
40,45	CLR: Undo T1 LSN 10, T1 End
50	update: T3 writes P1
60	update: T2 writes P5
	CRASH, RESTART
64,65,70	T2 abort, T3 abort, CLR: Undo T2 LSN 60
80,85	CLR: Undo T3 LSN 50, T3 end
	CRASH, RESTART
90, 95	CLR: Undo T2 LSN 20, T2 end

undonextLSN

Additional Crash Issues

Source: Ramakrishnan & Gehrke (partial)



Additional Crash Issues

- What happens if system crashes during Analysis?
During REDO?



Additional Crash Issues

- What happens if system crashes during Analysis?
During REDO?
- How do you limit the amount of work in REDO?
 - Flush asynchronously in the background.
 - Watch “hot spots”!

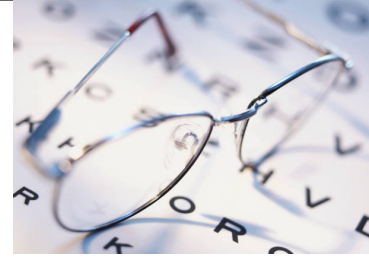


Additional Crash Issues

- What happens if system crashes during Analysis?
During REDO?
- How do you limit the amount of work in REDO?
 - Flush asynchronously in the background.
 - Watch “hot spots”!
- How do you limit the amount of work in UNDO?
 - Avoid long-running Xacts.



What should we learn today?



- Explain the concepts of volatile, nonvolatile, and stable storage as well as the main assumptions underlying database recovery
- Predict how force/no-force and steal/no-steal strategies for writes and buffer management influence the need for redo and undo
- Explain the notion of logging and the concept of write-ahead logging
- Predict what portions of the log and database are necessary for recovery based on the recovery equations
- Explain how write-ahead logging is achieved in the ARIES protocol
- Explain the functions of recovery metadata such as the transaction table and the dirty page table
- Interpret the contents of the log resulting from ARIES normal operation