



Reliability: Concepts + Replication

ACS, Yongluan Zhou

(many slides are made by Marcos Vaz Salles)

Highly-Available Systems

- Content distribution, web, media
 - E.g., YouTube
- Data Stores
 - E.g., Amazon Dynamo, Google F1
- Analytics
 - Long running jobs in Spark / MapReduce / Hadoop
 - Continuous stream processing, e.g., Storm / Samza / Kafka



Throughput Most Important Metric




Scale is the name of the game

- Volumes of information
 - TBs – PBs of raw data
- Variety of schemas and formats
 - Tens – hundreds of schemas
- Large data centers
 - Tens of thousands of machines



The Machinery



So, everything works?

- Assume a computer has probability of failure p
- If system needs N computers to work, what is probability of system working?
- Probability of one component working: $1-p$
- Probability of all components working: $(1-p)^N$
 - Assuming failures are independent!
 - Correlated failures are the reality – and make it even worse



Reliability Measures

- Mean Time to Failure (**MTTF**)
 - Mean Time to Repair (**MTTR**)
 - Mean Time Between Failures (**MTBF**)
 - **MTBF** = **MTTF** + **MTTR**
 - Availability = **MTTF** / **MTBF**
 - Downtime = (1 - Availability) = **MTTR** / **MTBF**
-
- Consider **N = 10,000** and for one computer,
MTTF = 30 years
-
1. How to estimate the value of MTTF for a system that has a long MTTF?
 2. How often do you estimate to see a computer failing in the above scenario?



Reliability & Availability

- Things will crash. Deal with it!
 - Assume you could start with super reliable servers (MTBF of 30 years)
 - Build computing system with 10 thousand of those
 - **Watch one failure per day**
 - Facebook* has to mitigate one datacenter outage every two weeks!
- Fault-tolerant software is inevitable
- Typical yearly flakiness metrics
 - 1-5% of your disk drives will die
 - Servers will crash at least twice (2-4% failure rate)

*Kaushik Veeraraghavan, Justin Meza, et al. Maelstrom: Mitigating Datacenter-level Disasters by Draining Interdependent Traffic Safely and Efficiently. In OSDI'2018.



The Joys of Real Hardware


Typical first year for a new cluster:

- ~0.5 **overheating** (power down most machines in <5 mins, ~1-2 days to recover)
- ~1 **PDU failure** (~500-1000 machines suddenly disappear, ~6 hours to come back)
- ~1 **rack-move** (plenty of warning, ~500-1000 machines powered down, ~6 hours)
- ~1 **network rewiring** (rolling ~5% of machines down over 2-day span)
- ~20 **rack failures** (40-80 machines instantly disappear, 1-6 hours to get back)
- ~5 **racks go wonky** (40-80 machines see 50% packetloss)
- ~8 **network maintenances** (4 might cause ~30-minute random connectivity losses)
- ~12 **router reloads** (takes out DNS and external vips for a couple minutes)
- ~3 **router failures** (have to immediately pull traffic for an hour)
- ~dozens of minor **30-second blips for dns**
- ~1000 **individual machine failures**
- ~thousands of **hard drive failures**
- **slow disks, bad memory, misconfigured machines, flaky machines, etc.**

Long distance links: **wild dogs, sharks, dead horses, drunken hunters, etc.**



Understanding Downtime Behavior Matters



Faults, Errors, and Failures

- Fault
 - Defect that has potential to cause problems
- Error
 - Wrong result caused by an active fault
- Failure
 - Unhandled error that causes interface to break its contract



Fault Tolerance

- Error detection
 - Use limited redundancy to verify correctness
 - Example: detect damaged frames in link layer
 - **Fail fast:** report error at interface
- Error containment
 - Limiting propagation of errors
 - Example: enforced modularity
 - **Fail stop:** immediately stop to prevent propagation
 - **Fail safe:** transform wrong values into conservative “acceptable” values, but limiting operation
 - **Fail soft:** continue with only a subset of functionality



Fault Tolerance

- Error masking
 - Ensure correct operation despite errors
 - Example: reliable transmission, process pairs
- We will focus on error masking
 - Main techniques



Replication


- **MAKE COPIES!!** 😊

- State-machine replication
- Asynchronous replication
 - Primary-Site
 - Peer-to-Peer
- Synchronous replication
 - Read-Any, Write-All
 - Quorums



Replicated
Interpreter

```
(loop (print (eval (read))))
```



Replicated
memory




- Techniques only good enough for a specific **failure model**
 - Nuclear holocaust
 - Component maliciously outputs random gibberish (**Byzantine**)
 - Components **crash** without telling you anything
 - Components are **fail-stop**

Asynchronous Replication




Asynchronous Replication



- Allows **WRITES** to return before all copies have been changed
 - **READs** nonetheless look at subset of copies
 - Users must be aware of which copy they are reading, and that copies may be out-of-sync for short periods of time.
- Two approaches: **Primary Site** and **Peer-to-Peer** replication
 - Difference lies in how many copies are “*updatable*” or “*master copies*”.




Primary Site Replication



- Exactly one copy is designated the **primary** or **master** copy. Replicas at other sites cannot be directly updated
 - The primary copy is **published**
 - Other sites **subscribe** to this copy; these are **secondary** copies
- Main issue: How are changes to the primary copy propagated to the secondary copies?
 - Done in two steps: First, **CAPTURE** changes made at primary; then **APPLY** these changes
 - Many possible implementations for **CAPTURE** and **APPLY**




Primary copy



- Writers lock & update primary copy and propagate the update to other copies
- Readers lock and access primary copy
- Widely adopted, e.g. many database systems

Peer-to-Peer Replication



- More than one of the copies of an object can be a master in this approach
 - Changes to a master copy must be propagated to other copies
 - If two master copies are changed in a conflicting manner, this must be resolved. (e.g., Site 1: Joe's age changed to 35; Site 2: to 36)
- Best used when conflicts do not arise
- **Examples**
 - Each master site owns a disjoint fragment of the data
 - Updating rights owned by one master at a time
 - Operations are associative-commutative



Eventual consistency

- If no new updates are made to an object, after some inconsistency window closes, all accesses will return the same “last” updated value
- **Prefix property:**
 - If Host 1 has seen write $w_{i,2}$: i^{th} write accepted by host 2
 - Then 1 has all writes $w_{j,2}$ (for $j < i$) accepted by 2 prior to $w_{i,2}$
- Assumption: write conflicts will be easy to resolve
 - Even easier if whole-“object” updates only

How do we know
the precedence
relationship?



Events and Histories


- Processes execute sequences of **events**
- Events can be of 3 types:
 - **local**, **send**, and **receive**
- The local history h_p of process p is the sequence of events executed by process



Ordering Events


- **Observation 1:**

- Events in a local history are totally ordered



- **Observation 2:**

- For every message m , $\text{send}(m)$ precedes $\text{receive}(m)$




Happens-Before (Lamport [1978])


- **Relative time? Define Happens-Before (\rightarrow) :**
 - On the same process: $a \rightarrow b$, if $\text{time}(a) < \text{time}(b)$
 - If p1 sends m to p2: $\text{send}(m) \rightarrow \text{receive}(m)$
 - Transitivity: If $a \rightarrow b$ and $b \rightarrow c$ then $a < c$
- **Lamport Algorithm establishes partial ordering:**
 - All processes use counter (clock) with initial value of 0
 - Counter incremented / assigned to each event as timestamp
 - A send (msg) event carries its timestamp
 - For receive (msg) event, counter is updated by $\max(\text{receiver-counter}, \text{message-timestamp}) + 1$



Events Occurring at Three Processes




Lamport Timestamps




Lamport Logical Time

- Fill in the missing values, A-F:




Lamport Logical Time


- Fill in the missing values, A-F:



Lamport Logical Time



Can we say: if $\text{timestamp}(e) < \text{timestamp}(f)$ then e precedes f ?



Logically concurrent events!

Source: Freedman



Vector Logical Clocks

- **With Lamport Logical Time**

- e precedes $f \Rightarrow \text{timestamp}(e) < \text{timestamp}(f)$, but
- $\text{timestamp}(e) < \text{timestamp}(f) \not\Rightarrow e$ precedes f



Vector Logical Clocks

- **With Lamport Logical Time**

- e precedes $f \Rightarrow \text{timestamp}(e) < \text{timestamp}(f)$, but
- $\text{timestamp}(e) < \text{timestamp}(f) \not\Rightarrow e$ precedes f


- Vector Logical time guarantees this:

- All hosts use a vector of counters (logical clocks),
- i th element is the clock value for host i , initially 0
- Each host i , increments the i th element of its vector upon an event, assigns the vector to the event.
- A `send(msg)` event carries vector timestamp
- For `receive(msg)` event,

$$V_{\text{receiver}}[j] = \begin{cases} \text{Max}(V_{\text{receiver}}[j], V_{\text{msg}}[j]), & \text{if } j \text{ is not self} \\ V_{\text{receiver}}[j] + 1 & \text{otherwise} \end{cases}$$




Vector Timestamps



Vector Logical Time

- Fill in the missing values, A and B:



$$V_{\text{receiver}}[j] = \begin{cases} \text{Max} (V_{\text{receiver}}[j], V_{\text{msg}}[j]), & \text{if } j \text{ is not self} \\ V_{\text{receiver}}[j] + 1 & \text{otherwise} \end{cases}$$

Source: Freedman



- Fill in the missing values, A and B:

Host 1: 1,0,0,0 (red), 2,0,0,0 (blue)

Host 2: 1,1,0,0 (blue), 1,2,0,0 (blue)

Host 3: 2,0,2,0 (blue), 1,2,0,0 (red)

Host 4: 2,0,2,1 (blue)

$$V_{\text{receiver}}[j] = \begin{cases} \text{Max } (V_{\text{receiver}}[j], V_{\text{msg}}[j]), & \text{if } j \text{ is not self} \\ V_{\text{receiver}}[j] + 1 & \text{otherwise} \end{cases}$$

Comparing Vector Timestamps

- $a = b$ if they agree at every element
 - $a < b$ if $a[i] \leq b[i]$ for every i , but $!(a = b)$
 - $a > b$ if $a[i] \geq b[i]$ for every i , but $!(a = b)$
 - $a \parallel b$ if $a[i] < b[i]$, $a[j] > b[j]$, for some i, j
(conflict!)
-
- If one history is prefix of other, then one vector timestamp $<$ other
 - If one history is not a prefix of the other, then (at least by example) VTs will not be comparable.



Eventual is not the only choice

- Host of other properties available
 - Beyond our scope!
- **Examples**
 - Strong consistency
 - Weak consistency
 - Causal consistency
 - Read-your-writes consistency
 - Session consistency
 - Monotonic read consistency
 - Monotonic write consistency
- See Werner Vogels' entry
http://www.allthingsdistributed.com/2007/12/eventually_consistent.html for informal overview, or a good distributed systems book for algorithms 😊



Synchronous Replication



Synchronous Replication



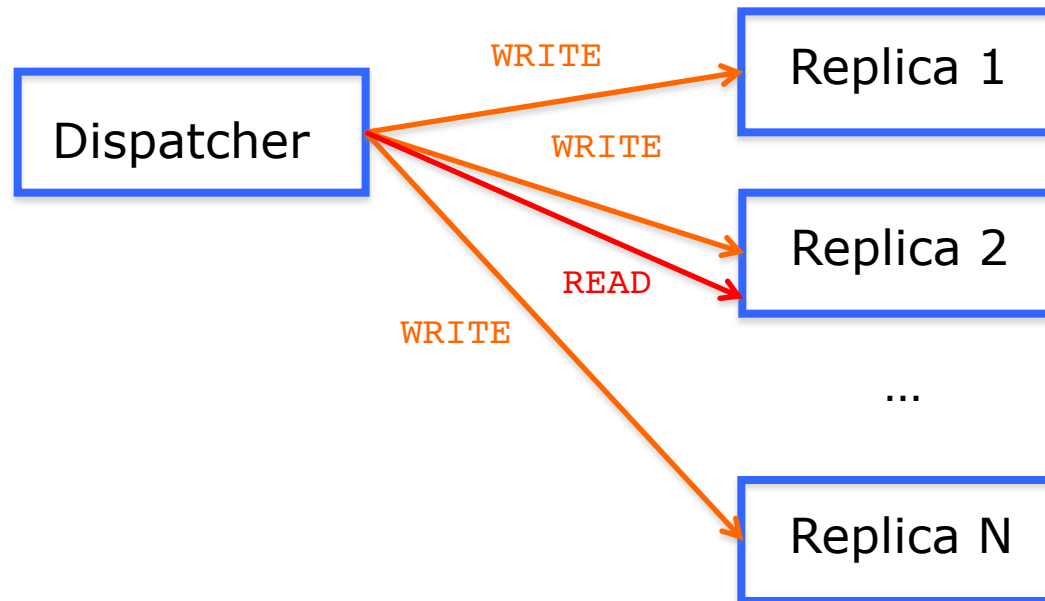
- Hide replication behind **READ**/**WRITE** memory abstraction
- Program operates against memory
- Memory makes sure **READS** and **WRITES** are **atomic**
 - **All-or-nothing:** either in all correct replicas or none
 - **Before-or-after:** Equivalent to a total order
- Memory replicates data for fault tolerance



Synchronous Replication



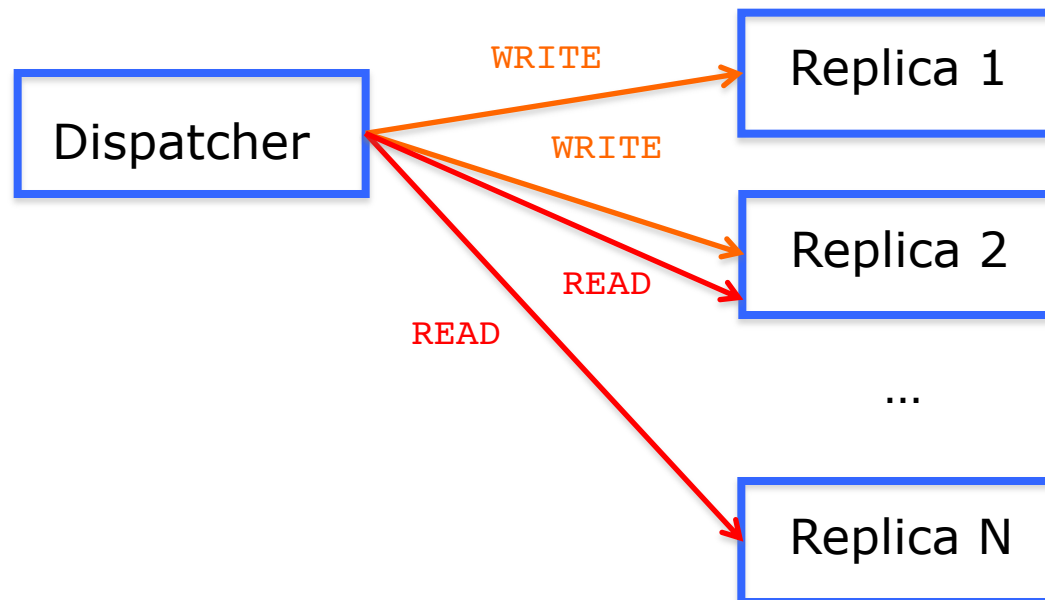
- **Read Any, Write-All**
 - For now assume we have a centralized Dispatcher → state-machine replication algorithms drop that assumption!
- WRITES synchronously sent everywhere
- But READS can be answered by any replica




Synchronous Replication



- Quorums
 - Read Quorum (Q_r) / Write Quorum (Q_w)
 - $Q_r + Q_w > N_{\text{replicas}}$
- Reads or writes only succeed if same response is given by respective quorum
 - Read any, Write all case is $Q_w = N_{\text{replicas}}$, $Q_r = 1$



What should we learn today?



- Explain and apply common fault-tolerance strategies such as error detection, containment, and masking
- Explain techniques for redundancy, such as n-version programming, error coding, duplicated components, replication
- Categorize main variants of replication techniques and implement simple replication protocols
- Explain the difficulties of guaranteeing atomicity in a replicated distributed system
- Discuss consistency properties in a replicated system and the notion of eventual consistency